



# CS 1550

Week 5 - Project 2

TA: Jinpeng Zhou      [jiz150@pitt.edu](mailto:jiz150@pitt.edu)

# CS 1550 – Project 2

---

- In Project 1, we have a modified Linux with our semaphore implementation, we also have a user-level program: trafficsim.
- In Project 2, you'll need to **write a specific user-level program**, which runs on the modified kernel from Project 1.
  - You can reuse your Project 1 Linux kernel
  - Or you can download a working kernel from Canvas

Build your user-level program on toth, then run it on qemu virtual machine.

# Project 2 - Museum Tour Simulation

---

- The specific user-level program is a museum tour simulation with some constraints, e.g.:
  - A visitor can tour the museum only if: (1) s/he has a ticket and (2) there's a guide on duty inside the museum. If no ticket is available, visitors will directly leave. If a visitor has a ticket but currently no guide on duty inside the museum, s/he will wait for a guide to be on duty.
  - A tour guide is on duty after arriving, and will be off duty after served 10 visitors or all tickets have been dispensed (i.e. there are no more visitors).
  - A tour guide cannot leave until s/he is off duty.
  - At most **two** tour guides can be in the museum at a time

Assume the museum knows there will be  $K$  guides and  $M$  visitors in a given day, i.e., the museum has  $K*10$  tickets available and it knows the number of sold tickets (the number of visitors to be served) is:  $\min(K*10, M)$

**Please refer to the project pdf for detailed constraints.**

# Project 2 - Safe Museum Tour Problem


---

- Your program **should always** satisfy those **constraints**, without causing any **deadlock**
- Your program takes several command-line arguments, such as number of visitors, number of tour guides, etc.

# Project 2 – Concurrent Execution

- Visitors and guides have different logics. Each visitor/guide is implemented as a child process.
- How can we create these child processes properly?

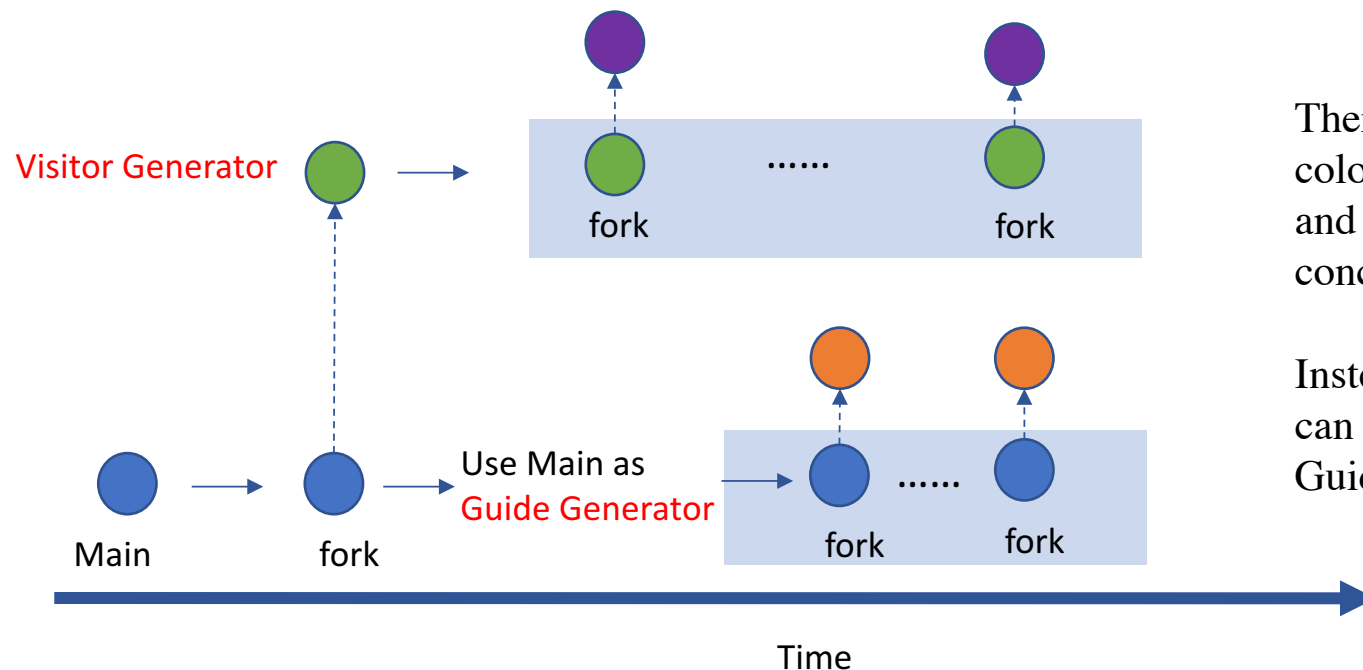
Intuitively, we may keep calling fork in the main(), e.g., as shown below. **However, with this method, the visitors and guides are not created concurrently**, due to a happen-before order: main creates all visitors before creating any guide.



# Project 2 – Concurrent Execution

- Visitors and guides have different logics. Each visitor/guide is implemented as a child process.
- How can we create these child processes properly?

In order to concurrently create visitors/guides, we introduce **generator process**, e.g:




There's no synchronization between the blue-colored rectangles, so the purple nodes(visitors) and orange nodes(guides) can be forked concurrently (without any happen-before order)

Instead of using Main as Guide Generator, you can also explicitly create a new child process as Guide Generator.

# Process hierarchy

---

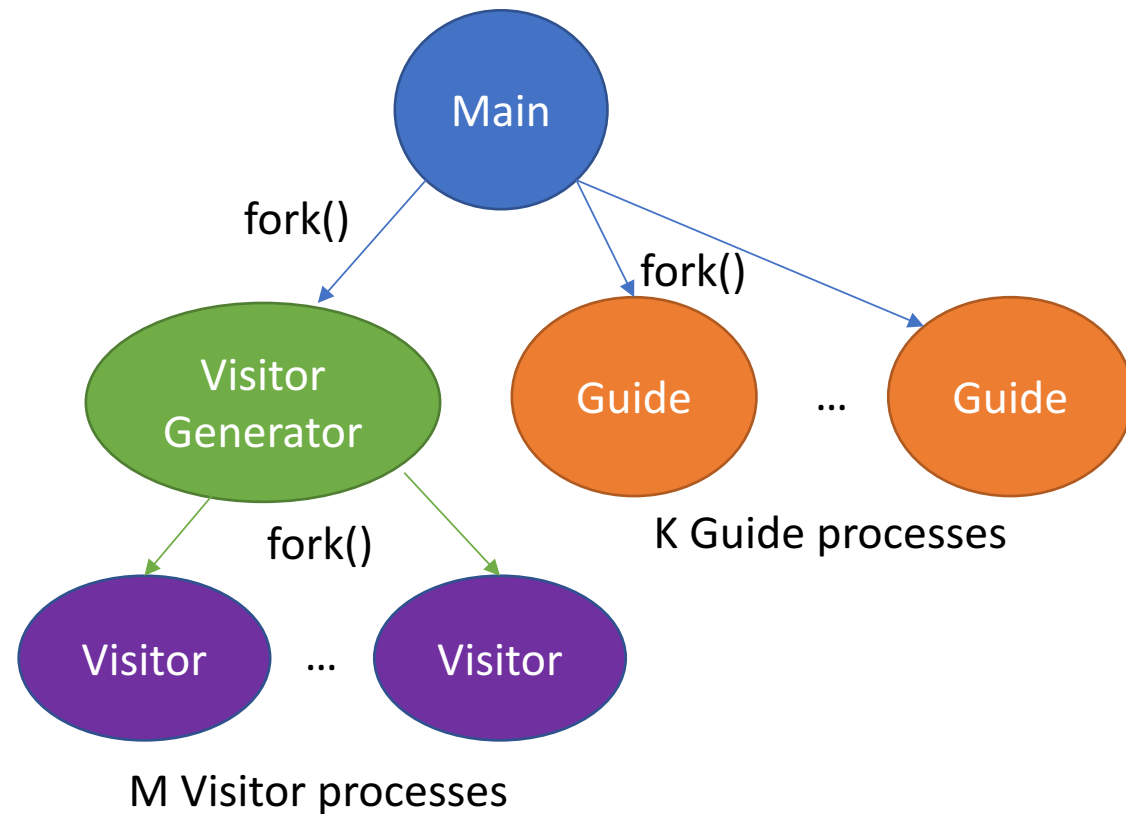
In our example, main process creates the visitor generator process.



# Process hierarchy

---

In our example, main process creates the visitor generator process.  
Then main process works as the guide generator process, and generate all guides.

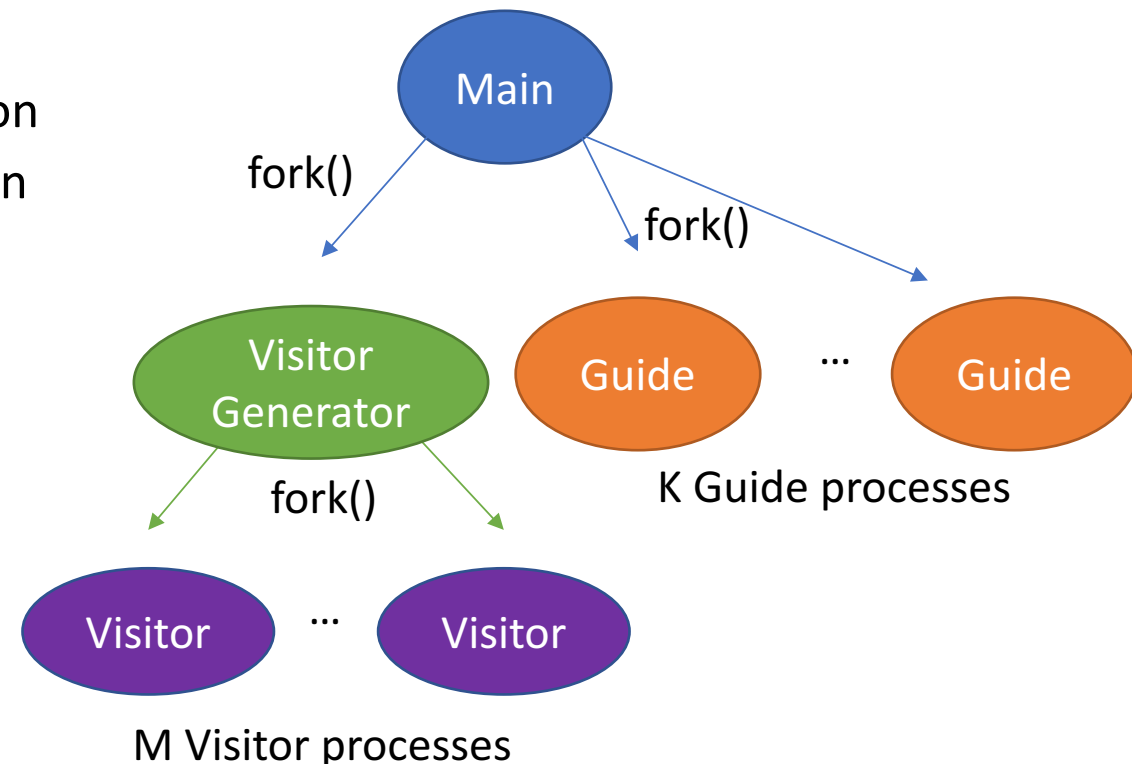




# wait() and exit()

---

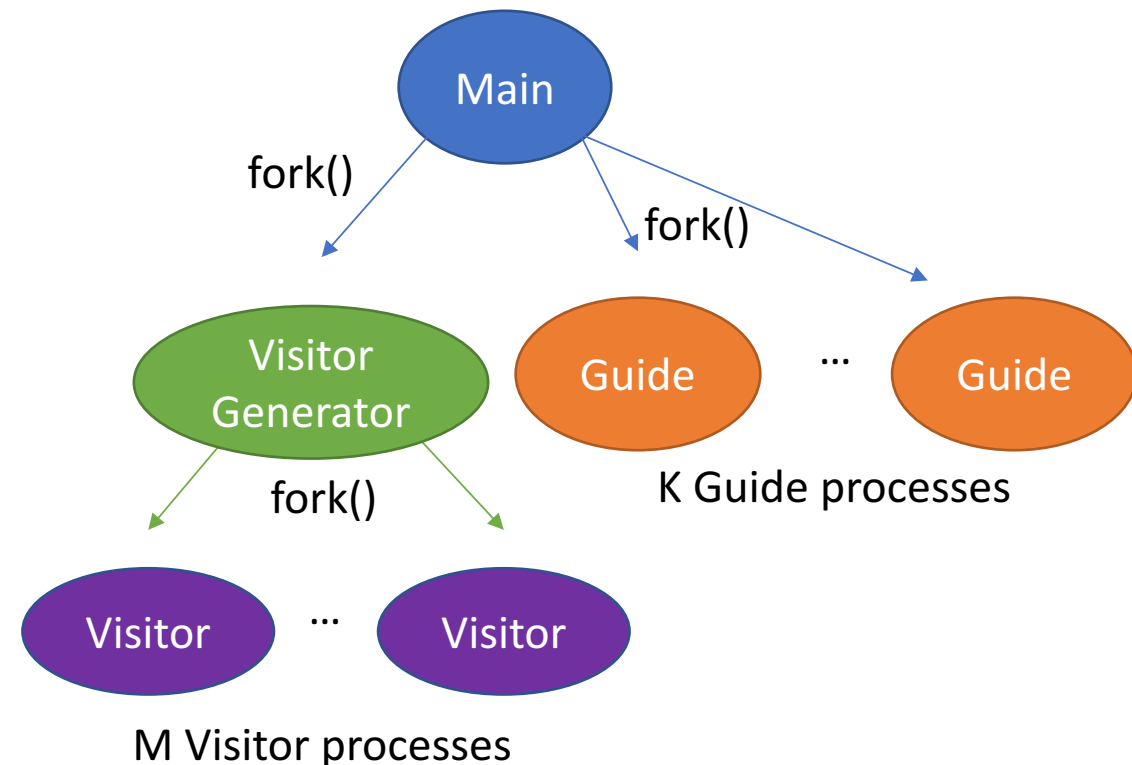
- A parent process must call “wait()”, waiting for its child to terminate, such that the resource occupied by the child process can be released. Otherwise:
  - If child terminates before parent, it becomes a “zombie”, still occupying resource after termination
  - The parent may terminate while child is running. In this case, the child becomes an “orphan”, and will be adopted by the Init process (the ancestor of all processes)



# wait() and exit()

---

- In our example:
  - Main is the parent of visitor generator
    - Main should call one wait() for visitor generator
  - Main is the parent of all guides
    - Main should call K wait() for K guides
  - Visitor generator is the parent of all visitors
    - Visitor generator should call M wait() for M visitors



# Simulating visitors' & guides' arrival

- Here's one possible implementation:

```
Main process: // herein we ignore cases such as failed fork for simplicity
```

```
/* Main process creates the visitor generator process */
```

```
pid = fork();
```

```
if pid is 0:
```

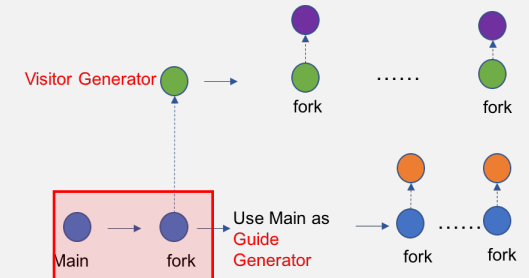
```
    /* Visitor generator process */
```

```
else if pid > 0:
```

```
    /* Main process (reused as Guide generator process) */
```

```
    // A straightforward way is to reuse the main process as guide generator,
```

```
    // or you can do a fork here to actually create the guide generator
```



# Simulating visitors' & guides' arrival

- Visitor generator process (guide generator process has a similar logic)

```
Visitor generator process: // herein we ignore cases such as failed fork for simplicity
```

```
/* Creates M visitors */
```

```
loop with M iterations:
```

```
    pid = fork();
```

```
    if pid is 0:
```

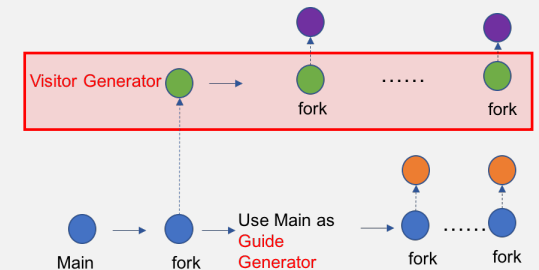
```
        /* Visitor process: implement your visitor tour logic here */
```

```
        . . .
```

```
    else:
```

```
        /* Generator process: proceed to generate next visitor */
```

```
        // Implement the probabilistic delay before generating next visitor
```



# Simulating visitors' & guides' arrival

---

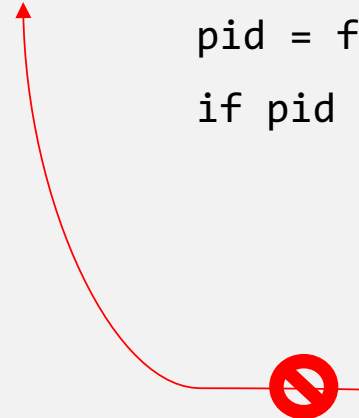
- Probabilistic delay example:

```
srand(seed);    // set seed for the generator
. . .
int value = rand() % 100 + 1;    //random number between 1 and 100
if (value > 30) {                // 70% chance to delay; 30% chance to create next visitor immediately
    sleep(delay_seconds);
}
```

# Simulating visitors' & guides' arrival

- Visitor process should exit after the visitor leaves the museum.

```
Visitor generator process: // herein we ignore cases such as failed fork for simplicity
/* Creates M visitors */
loop with M iterations:
    pid = fork();
    if pid is 0:
        /* Visitor process: implement your visitor tour logic here */
        . . .
        /* This visitor should not go back to loop after he finishes tour.
           Otherwise he will be creating visitors */
        . . .
    else:
        /* Generator process: proceed to generate next visitor */
```



# Share data among processes

---

- To synchronize visitors and guides, we'll need to check the current status, such as the the current number of visitors/guides inside museum, etc.
- These data should be shared among different processes.
  - Using global variable, local variable, or heap variable won't work, as a child process has its own copies of global/stack/heap regions:

```
pid = fork();  
/* child and parent both have a variable named "pid", but in different physical mem addr */  
if pid is 0:  
    /* child: "pid" equals 0 */  
else if pid > 0:  
    /* parent: "pid" equals the actual process_id of child */
```

# Share data among processes

---

- Use `mmap()` to allocate shared memory by using **MAP\_SHARED** flag

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int protection, int flags, int fd, off_t offset);
```

- Assume we'd like to share two 32-bit integers ( $2 * 4\text{Bytes} = \mathbf{8\ bytes}$  in total):

```
void *ptr = mmap(NULL, 8, PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS, 0, 0); // assume mmap returns 0x100
```

```
/* the first 4 bytes (0x100~0x103) is the 1st integer, the next 4 bytes (0x104~0x107) is the 2nd integer */
```

```
int *first = (int*) ptr; // set the first int pointer to the start of allocated memory, so the value of first is 0x100
```

```
int *second = first + 1; // adding 1 to an int pointer means adding sizeof(int) bytes, so the value of second is 0x104
```

Alternatively, you can define a struct that contains two ints, and use `mmap` to allocate mem for a struct object.

More info: <http://man7.org/linux/man-pages/man2/mmap.2.html>



# Print messages

---

- You'll need to print specific messages, e.g.:
  - When a tour arrives you should print
    - **Tour guide %d arrives at time d%**
  - When a visitor arrives you print
    - **Visitor %d arrives at time %d**
- Use these calls to print messages in multi-process programs:

```
printf(<print content here>);  
fflush(stdout);
```

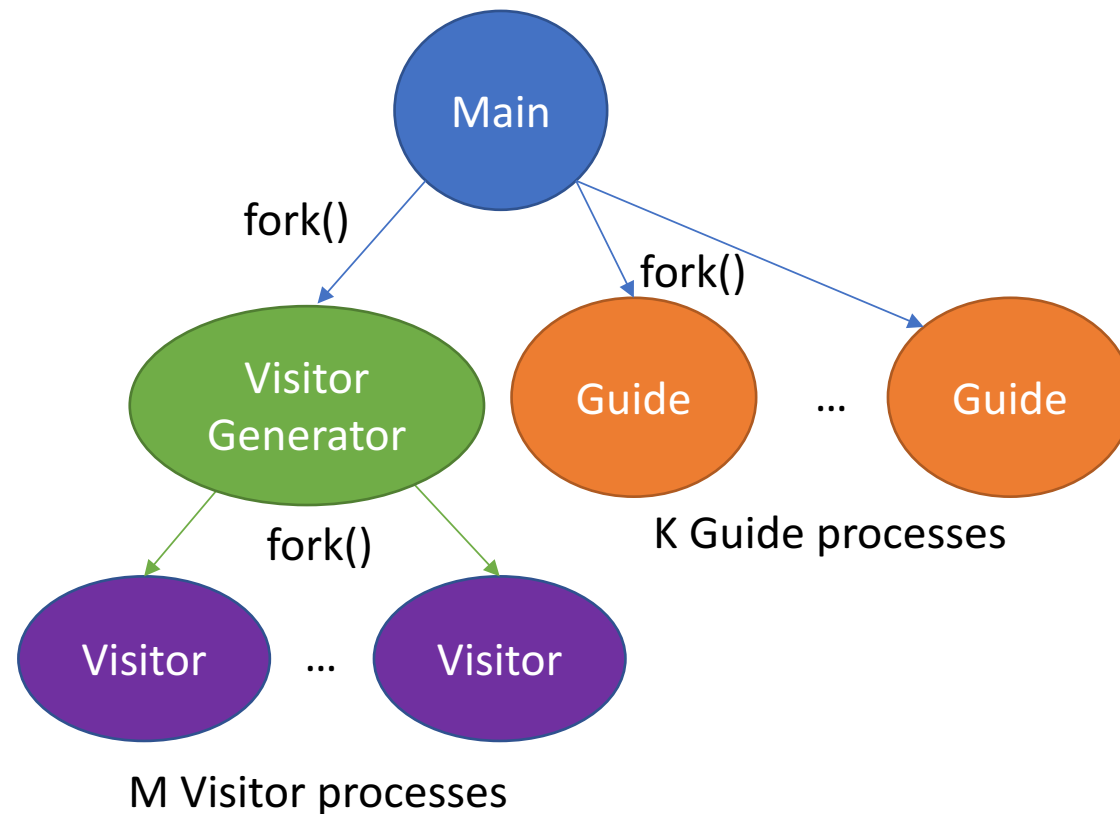
OR

```
fprintf(stderr,<print content here>);
```

# Tips

---

- First, try to fork()/wait() processes correctly for M visitors and K guides, without worrying about the detailed sync logic in visitor/guide.



# Tips

---

- Second, use mmap to create some shared variables
  - It should be done before creating processes.

```
void *ptr = mmap(NULL, 8, PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS, 0, 0);  
int *first = (int*) ptr;  
int *second = first + 1;
```

# Tips

---

- Then, think about detailed implementation of visitor and guide. We'll further discuss this next week.

## Visitor generator process:

```
/* Creates M visitors */
loop with M iterations:
    pid = fork();
    if pid is 0:
        /* Visitor process: implement your visitor tour logic here */
        . . . // synchronize, print, exit
    else:
        /* Generator process: proceed to generate next visitor */
```