



CS 1550

Week 8 - Project 2 Q/A

TA: Jinpeng Zhou jiz150@pitt.edu

Project2 New Deadline

- Due: Oct 23, 11:59pm
- Late: Oct 25, 11:59pm

Using provided kernel

- Kernel files are on thoth: /u/OSLab/original
 - **System.map**
 - **bzImage**
 - **linux/unistd.h**
 - **vmlinux**
- Download System.map and bzImage to vm, cp them to /boot, run lilo:
 - You only need to do this once, and make sure to reboot your vm with “devel”

Building museumsim

- The museumsim.c should include the unistd.h which defines the new syscall macros `__NR_cs1550_XXX`

```
#include <linux/unistd.h>
```

```
int main () {
```

```
}
```

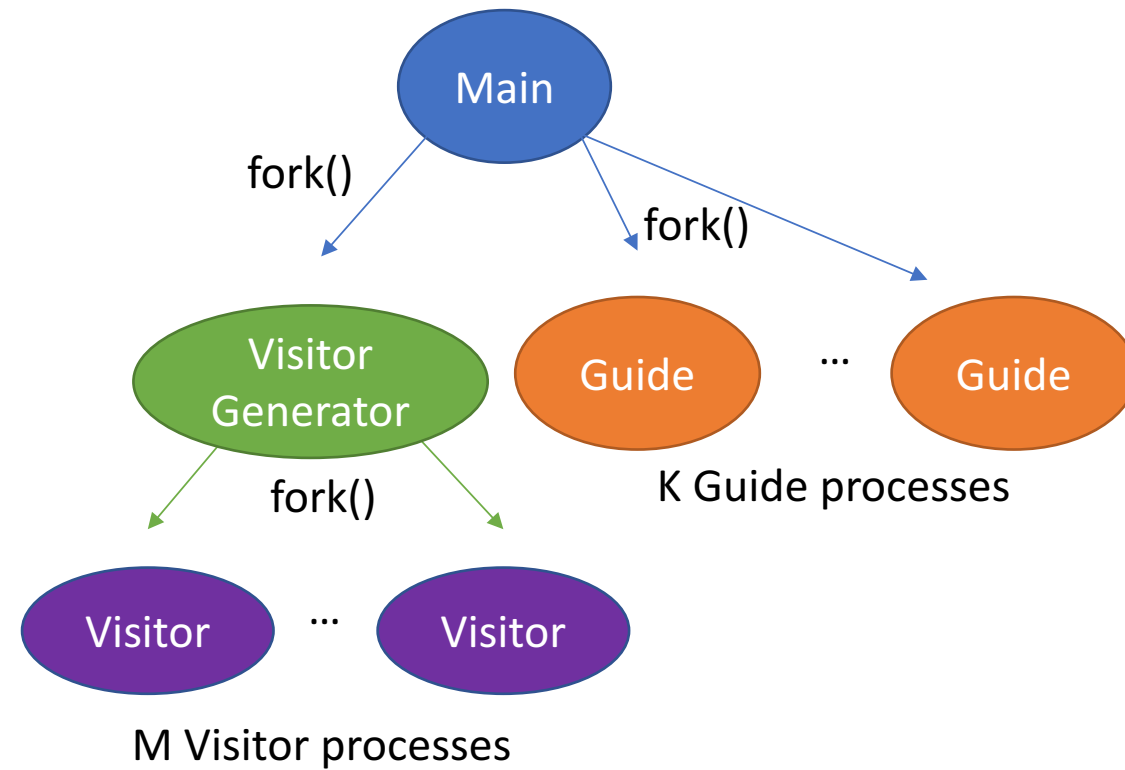
- We need to tell the compiler where it can find the `linux/unistd.h`. So, we'll use the following gcc command to build your program on thoth:

```
gcc -g -m32 -o museumsim -I /u/OSLab/original museumsim.c
```



Compiler will find linux/unistd.h inside /u/OSLab/original

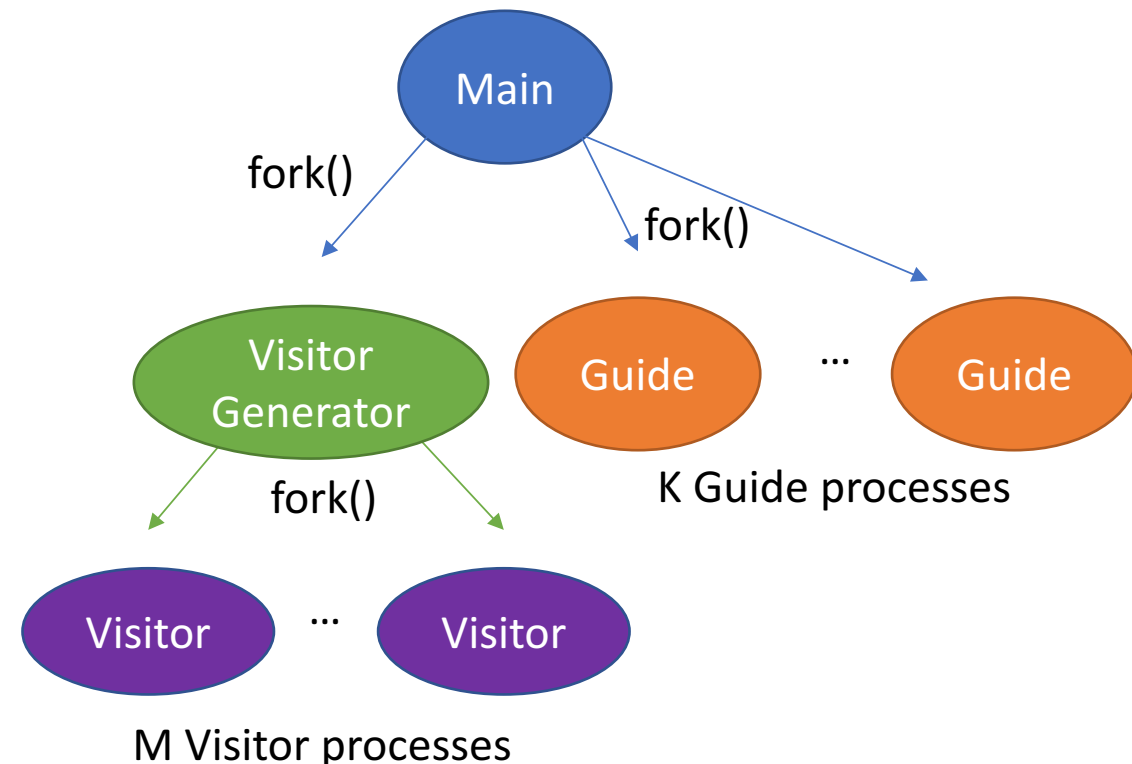
Creating visitors and guides



Managing processes

- A parent process must call “wait()” to wait for the child’s termination, such that the resource occupied by the child can be released
- In our example:
 - Main is the parent of visitor generator
 - Main should call **one wait() for visitor generator**
 - Main is the parent of all guides
 - Main should call **K wait() for K guides**
 - Visitor generator is the parent of all visitors
 - Visitor generator should call **M wait() for M visitors**

Make sure a parent with X children calls wait() X times.



Main process: // herein we ignore cases such as failed fork for simplicity

/* Main process */

pid = fork(); // Create visitor generator

if pid is 0:

/* Visitor generator process */

loop with M iterations: // Create M visitors

pid = fork();

if pid is 0:

/* Visitor process */

Visitor arrives, tours, leaves

Visitor process **exits**

else:

/* Visitor Generator process */

Probabilistic delay before generating next visitor

After Visitor Generator finishes all creation, it calls M wait() for M visitors

Visitor Generator process **exists**

else if pid > 0:

/* Main process (reused as Guide generator process) */

Similar logic as visitor generator

Main process calls one wait for visitor generator

return 0;

Sharing memory among processes

```
struct shared_mem {
    int data;
};

struct shared_mem * shared = NULL;

Main() {
    // Allocation
    shared = (struct shared_mem*) mmap (NULL, sizeof(struct shared_mem), PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS, 0, 0);

    // Initialization
    shared->data = 1;

    ...

    if (pid == 0){ // child process
        ...
        shared->data = 2;
        ...
    } else { // parent process
        ...
        shared->data = 3;
        ...
    }
}
```


Using semaphores

```
#include <linux/unistd.h>

long create(int value, char name[32], char key[32]) {
    return syscall(__NR_cs1550_create, value, name, key);
}

long open(char name[32], char key[32]) {
    return syscall(__NR_cs1550_open, name, key);
}

long down(long sem_id) {
    return syscall(__NR_cs1550_down, sem_id);
}

long up(long sem_id) {
    return syscall(__NR_cs1550_up, sem_id);
}

long close(long sem_id) {
    return syscall(__NR_cs1550_close, sem_id);
}
```

To achieve mutual exclusive execution:

```
long sem_mutex = create (1, name, key);
```

```
down (sem_mutex);
// access shared data
up (sem_mutex);
```

To represent a type of resource:

```
long sem_res = create (N, name, key); // N available initially
```

Please refer to trafficsim.c and trafficsim-mutex.c for more examples