



# CS 1550

Week 10 – Lab 4

Teaching Assistant

Henrique Potter

# CS 1550 – Lab 4

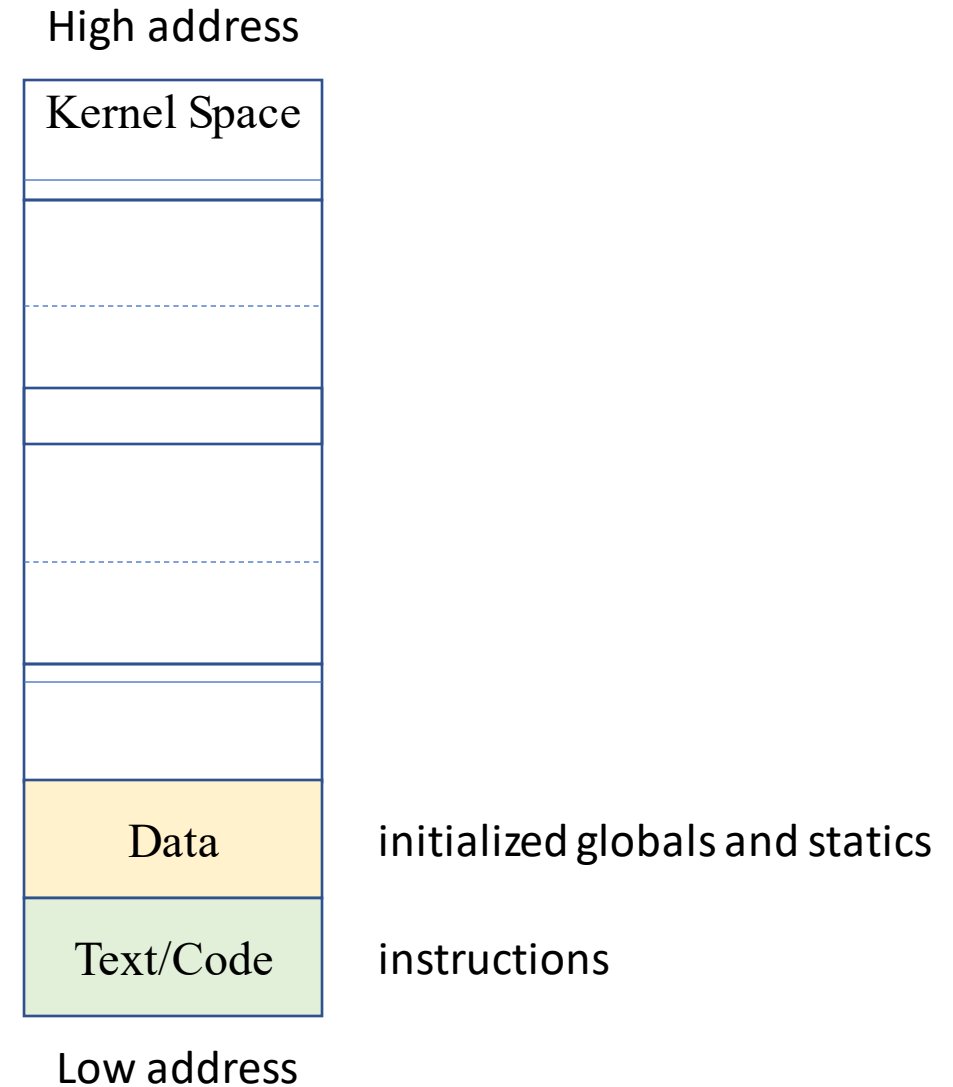
---

- Implementing lazy page allocation in xv6

# Memory layout

---

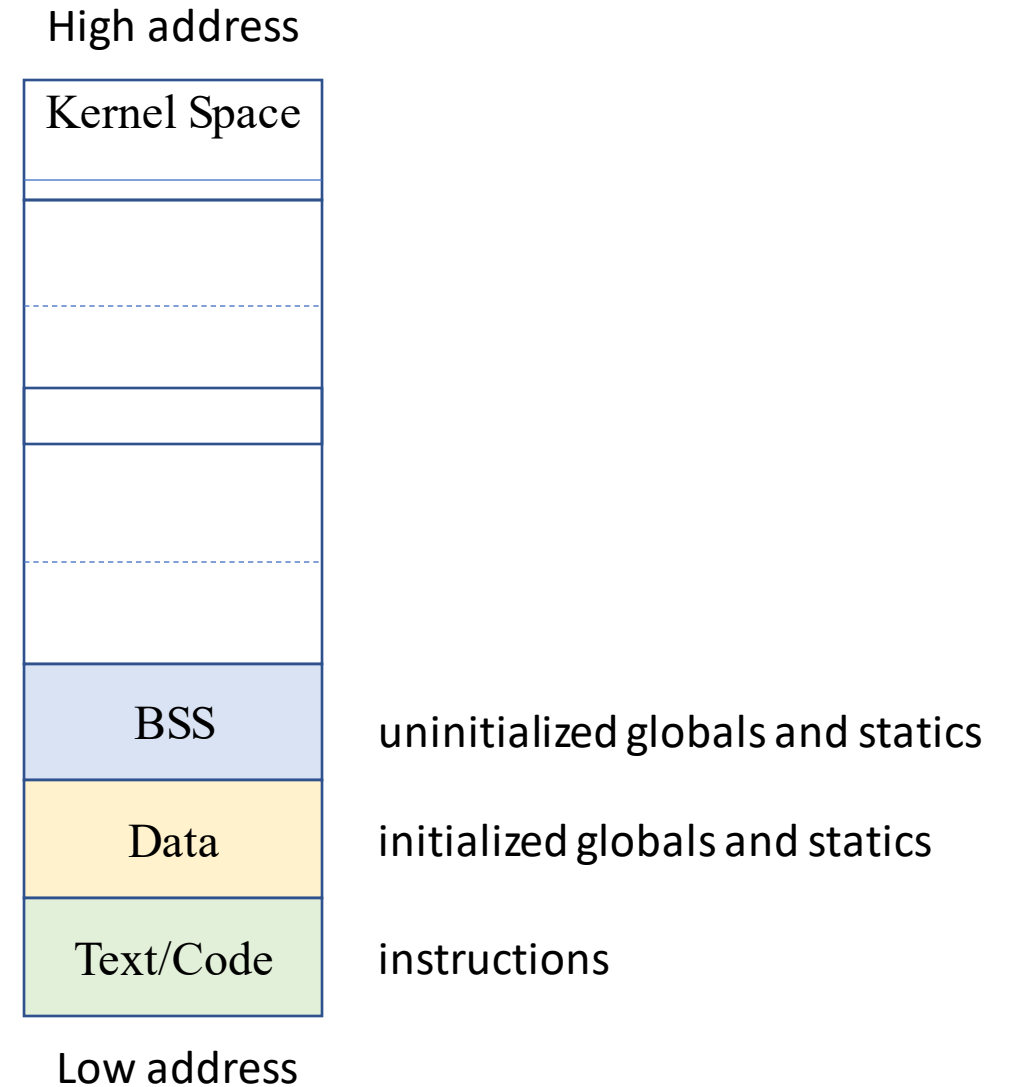
```
int t = 0; // Data
...
int main() {
    ...
}
```



# Memory layout

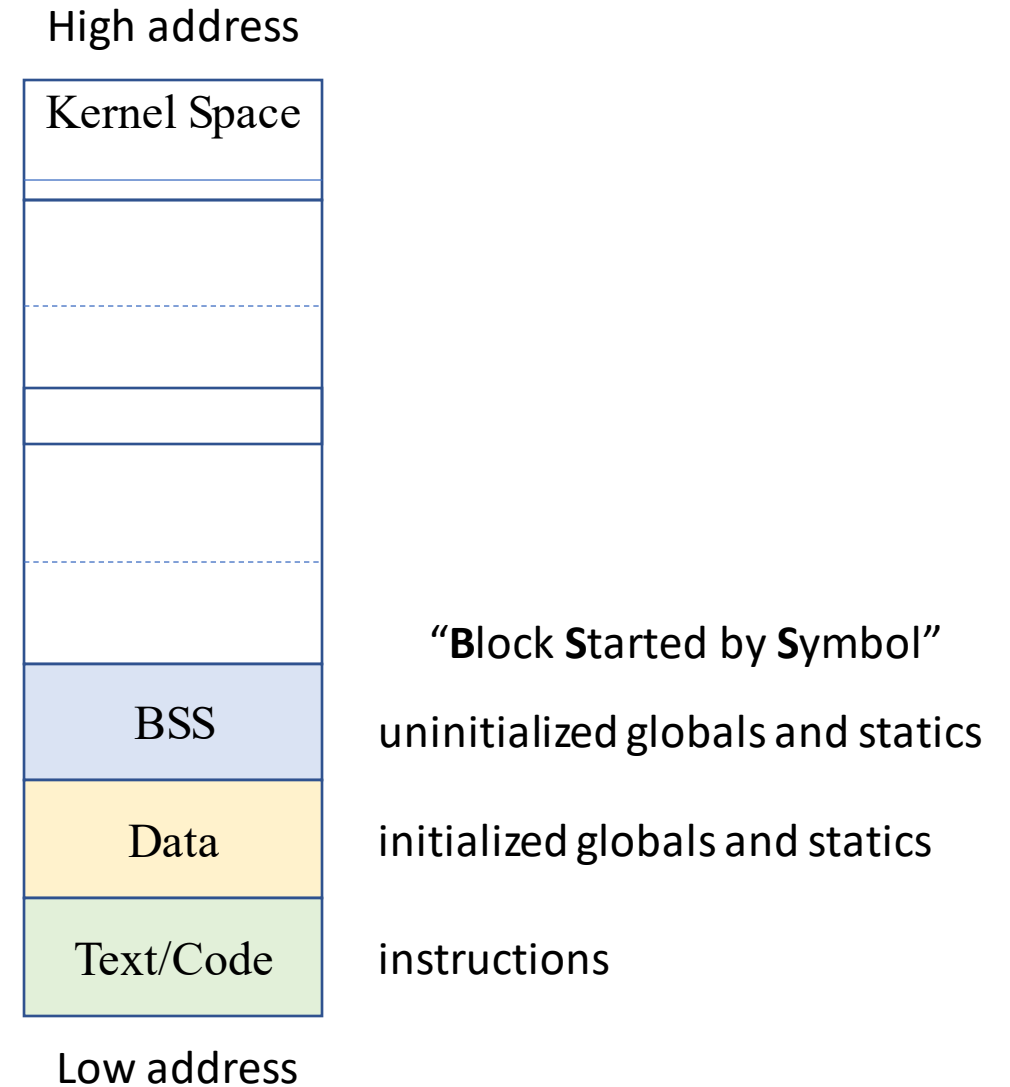
---

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    static int j;    // BSS
}
```



# Memory layout

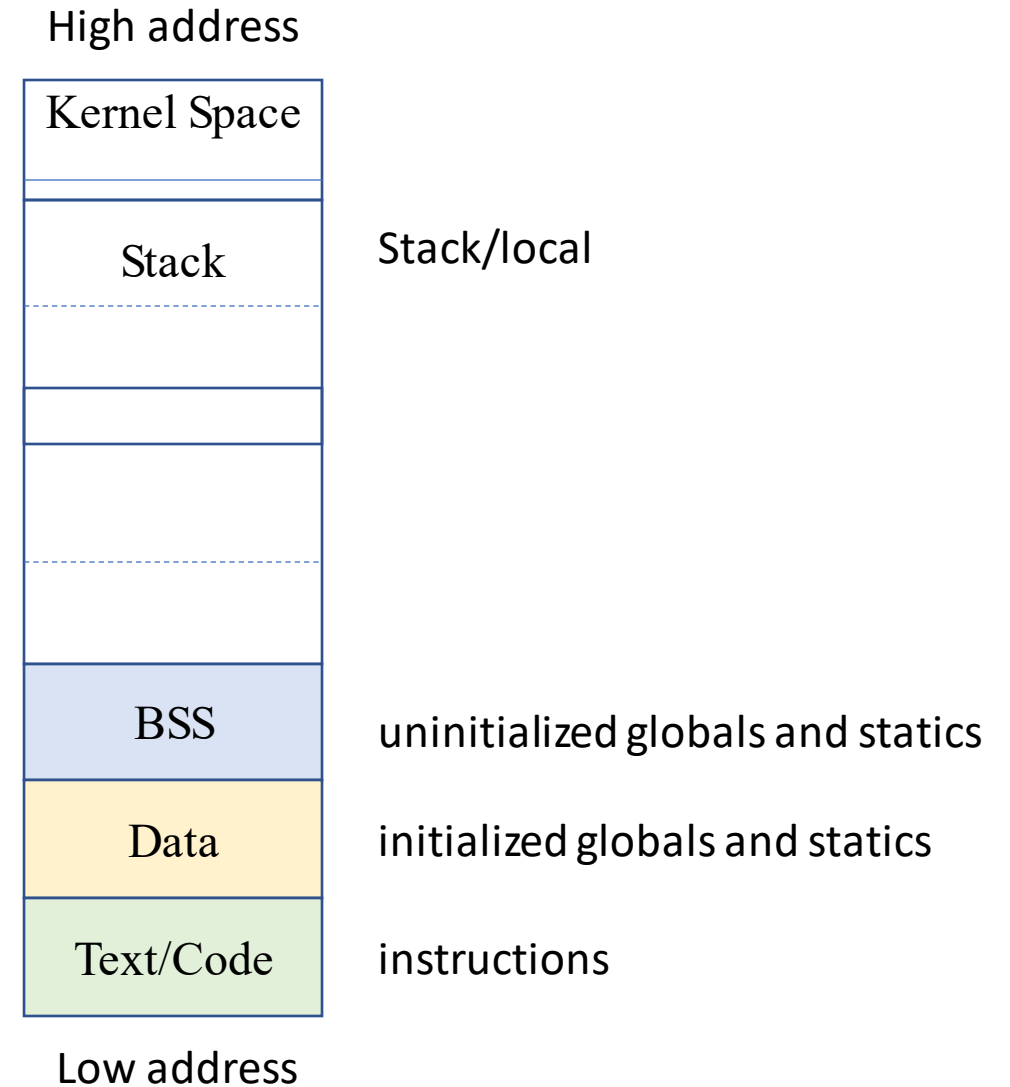
```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    static int j;    // BSS
}
```



# Memory layout

---

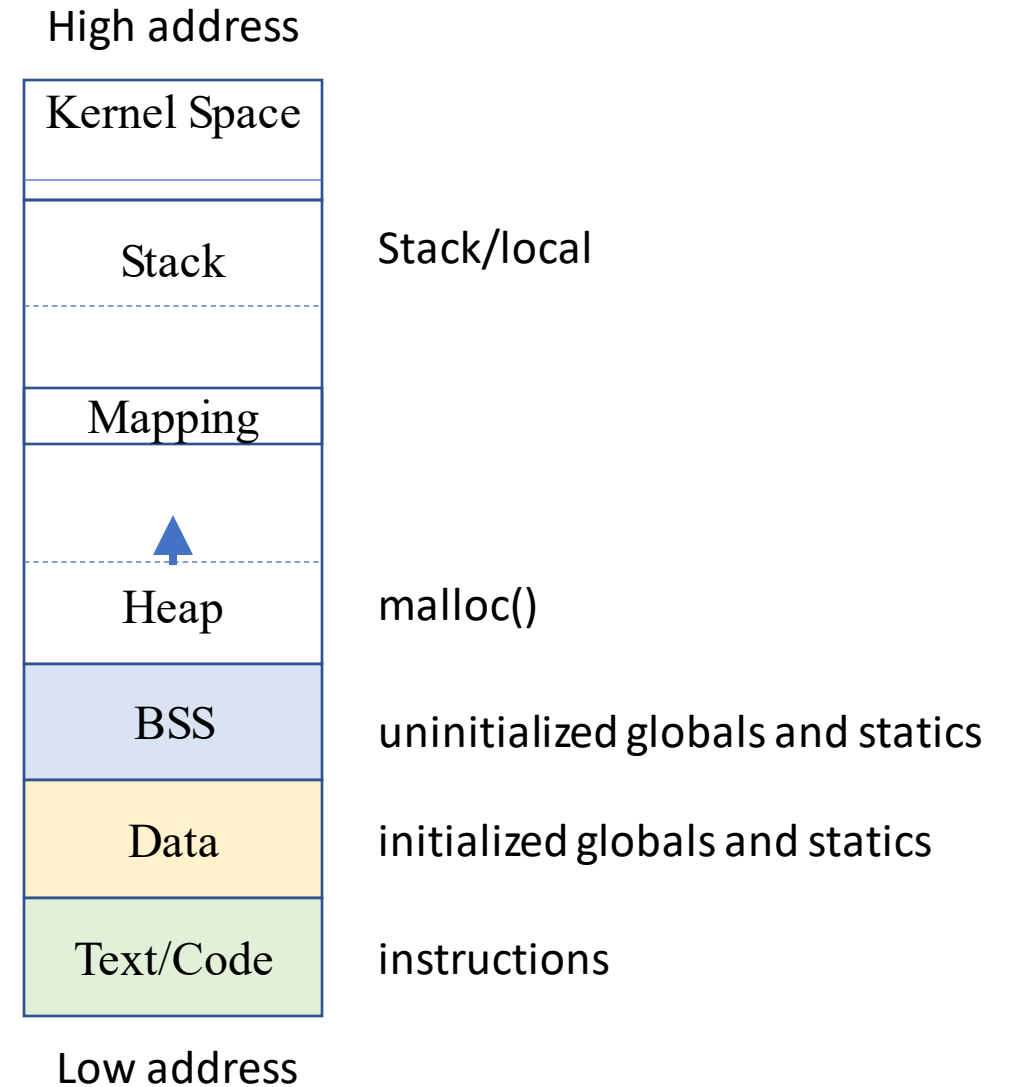
```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS
}
```



# Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS

    // ptr: Stack
    // 4B pointed by ptr: Heap
    char * ptr = (char*)malloc(4);
}
```




# Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS

    // ptr: Stack
    // 4B pointed by ptr: Heap
    char * ptr = (char*)malloc(4);

    // mptr: Stack
    // 4K pointed by mptr: memory Mapping
    char * mptr = (char*)mmap(...,4096,...);
    ...
}
```





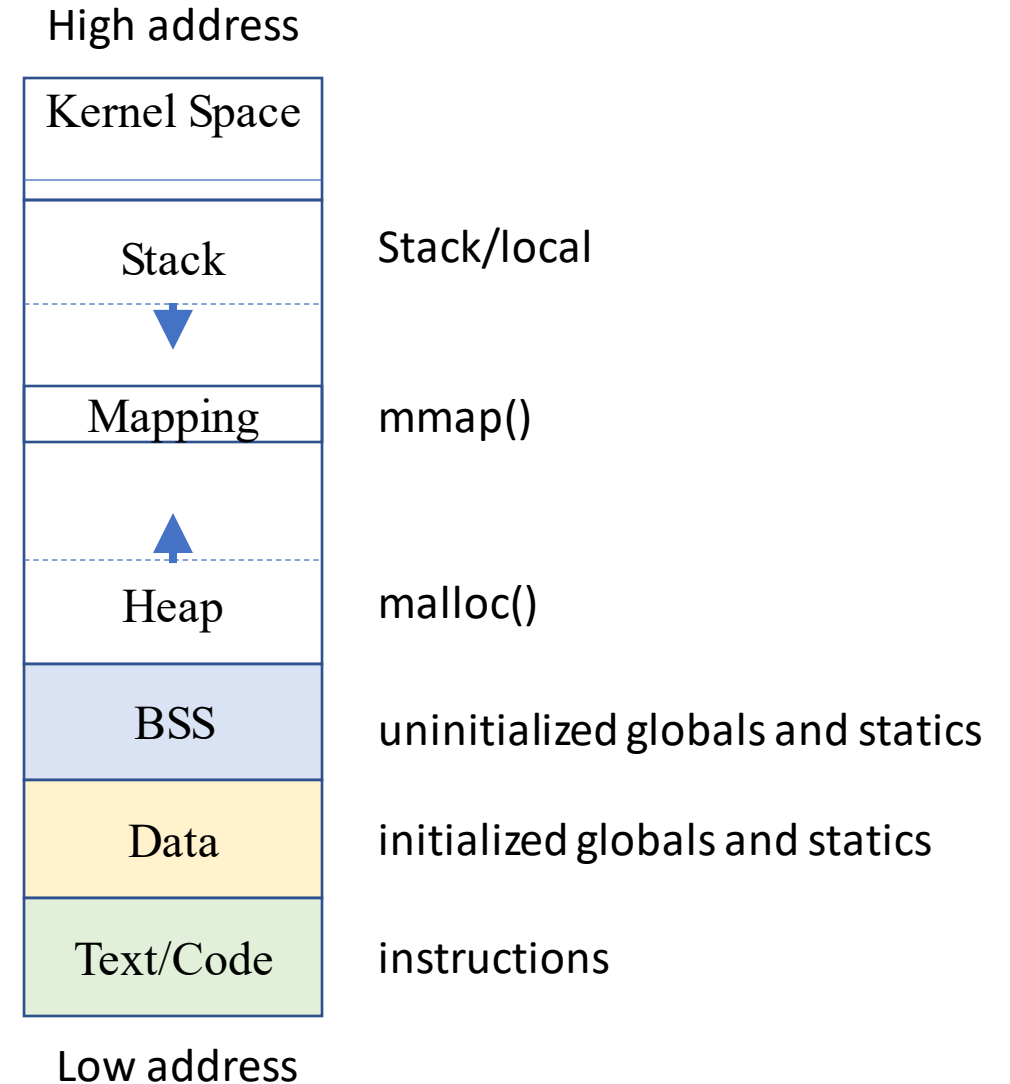
# Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS

    // ptr: Stack
    // 4B pointed by ptr: Heap
    char * ptr = (char*)malloc(4);

    // mptr: Stack
    // 4K pointed by mptr: memory Mapping
    char * mptr = (char*)mmap(...,4096,...);
    ...
}
```

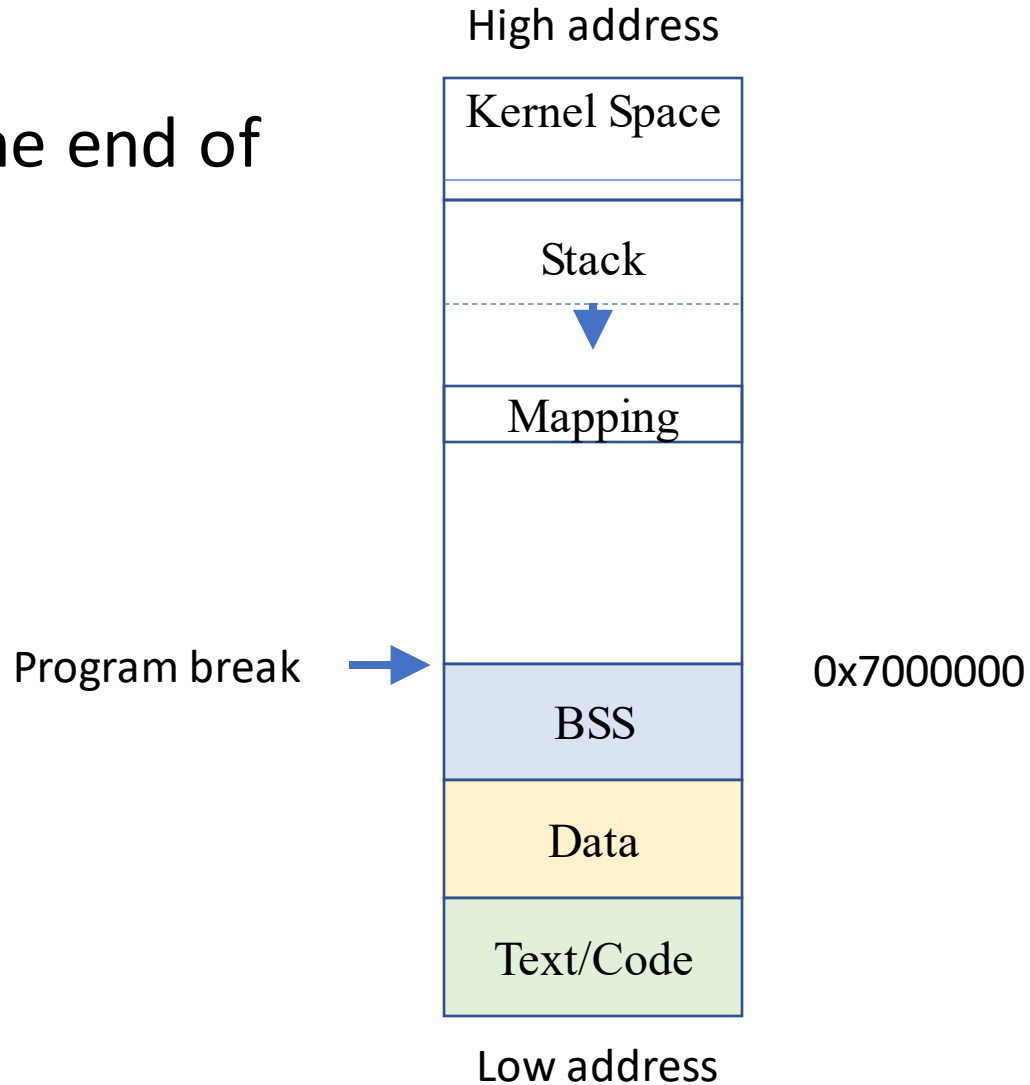
**xv6 do not have malloc!**



# Program Break

---

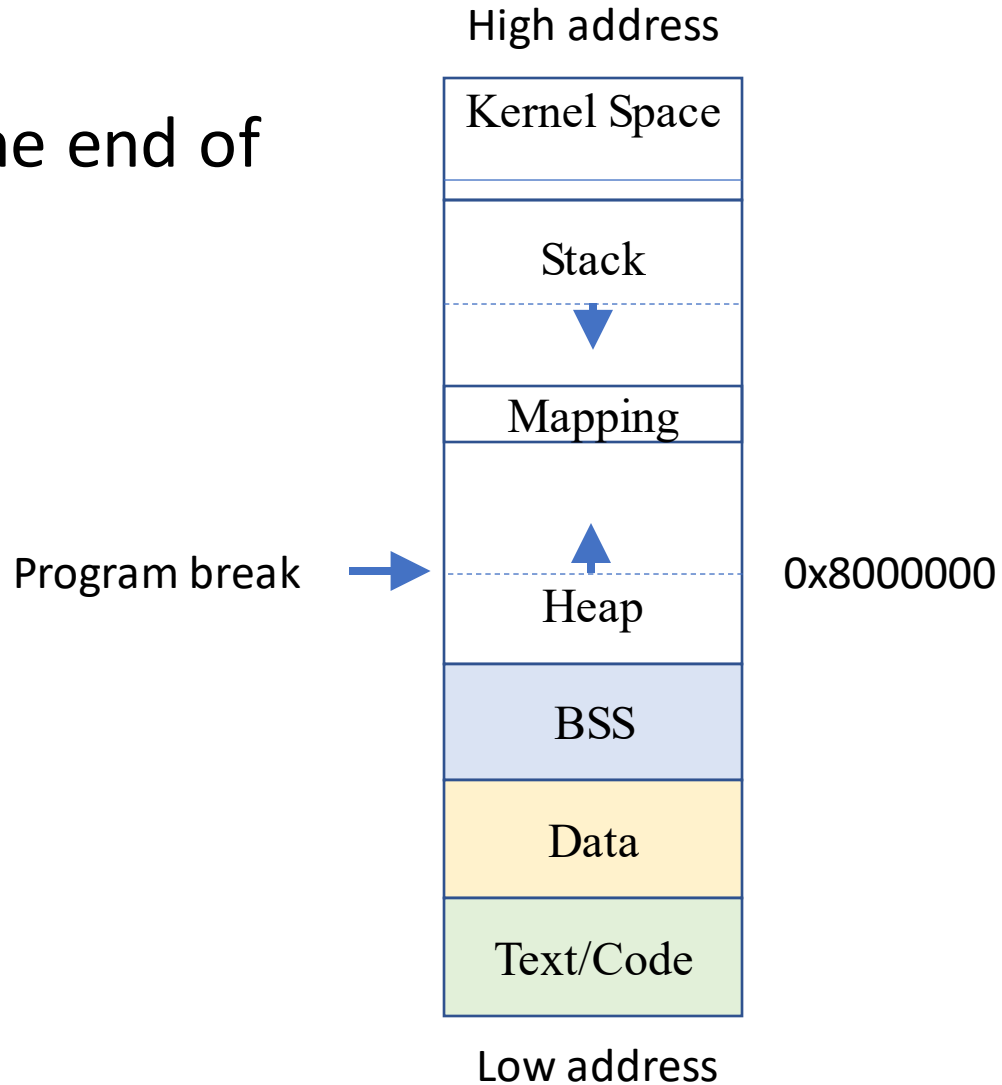
- Program break marks the end of the uninitialized data



# Program Break

---

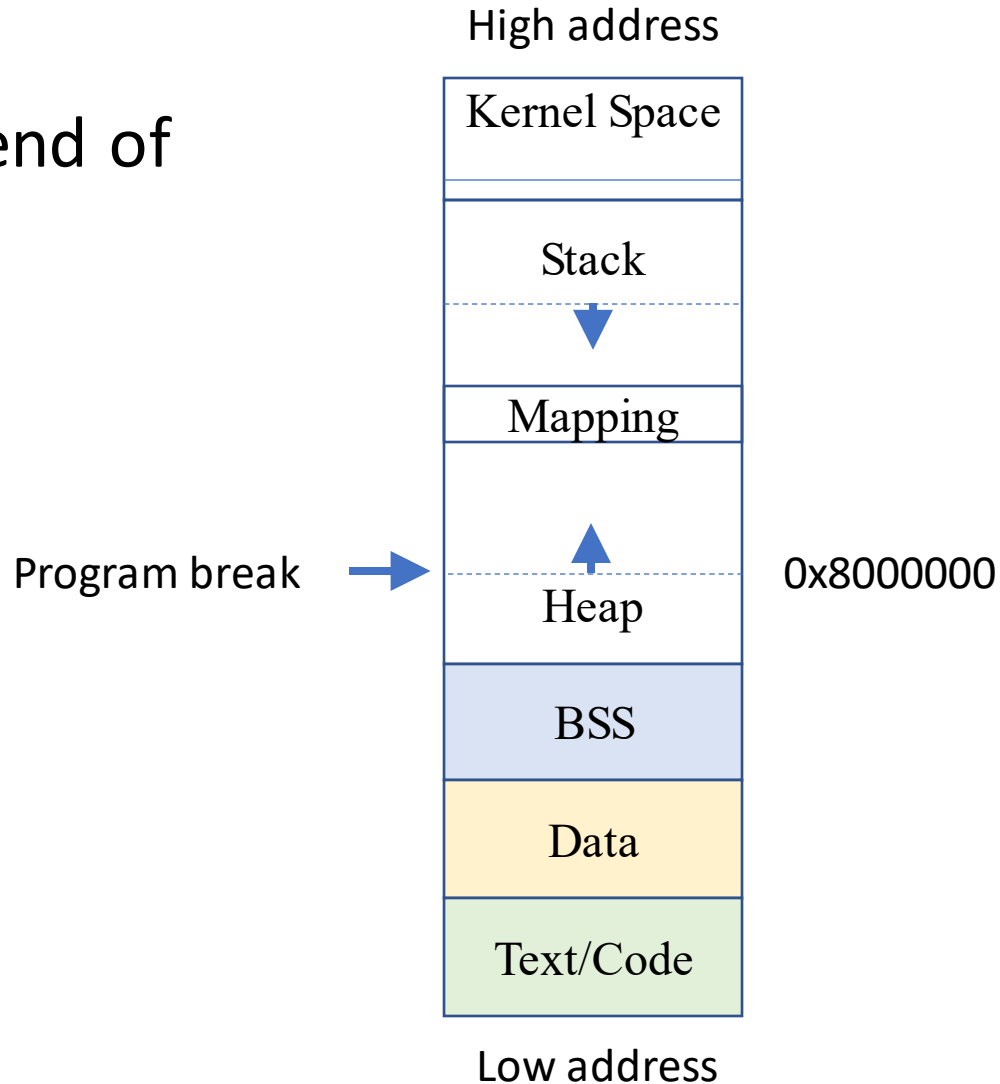
- Program break marks the end of the uninitialized data



# Program break: The syscall brk

- Sbrk adds a size to the end of cur\_brk

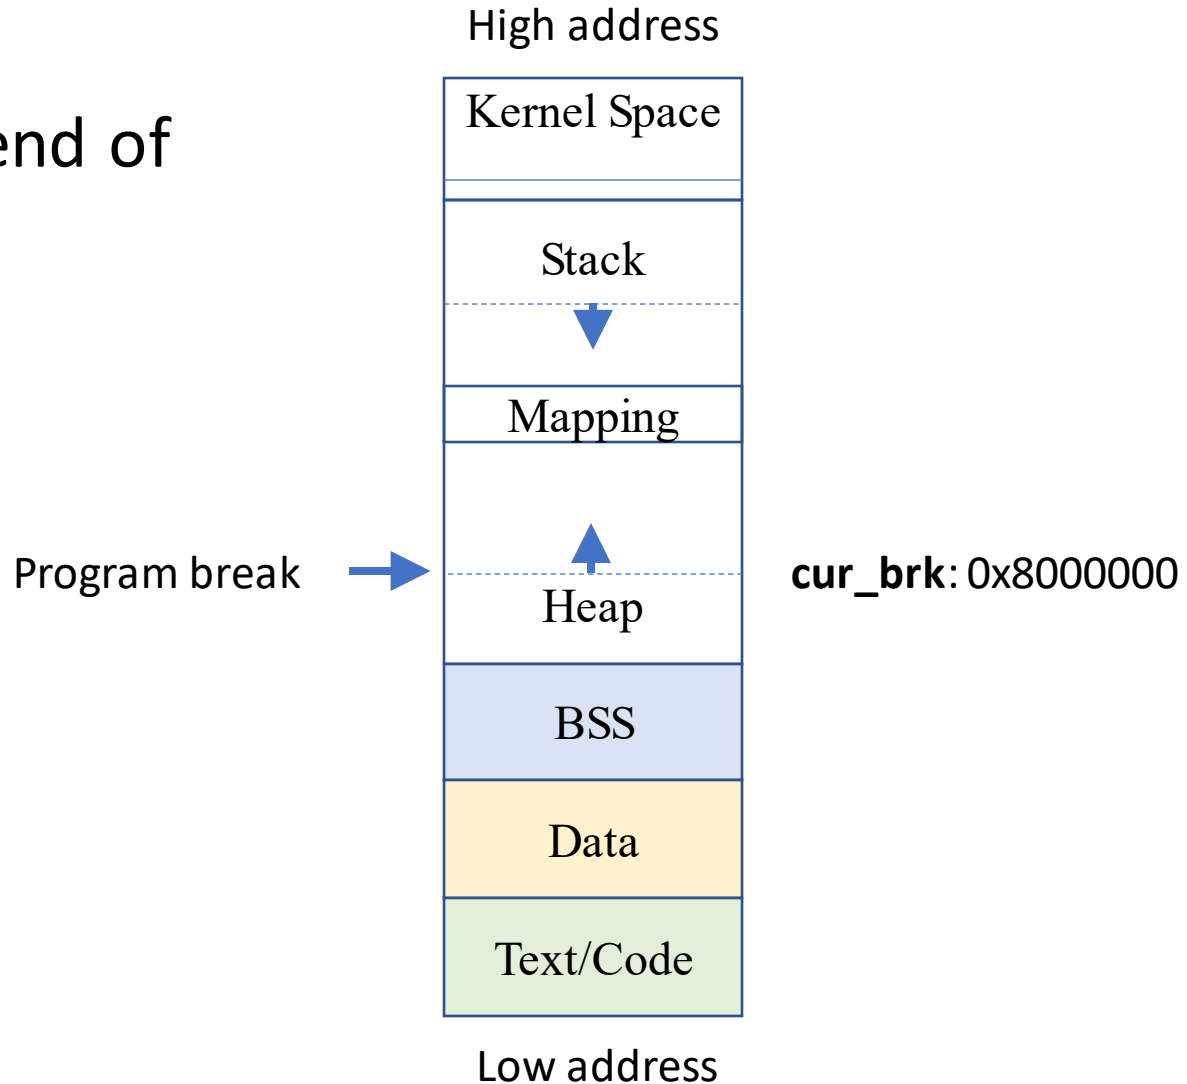
```
void *cur_brk = sbrk(0);
```



# Program break: The syscall brk

- Sbrk adds a size to the end of cur\_brk

```
void *cur_brk = sbrk(0);
```



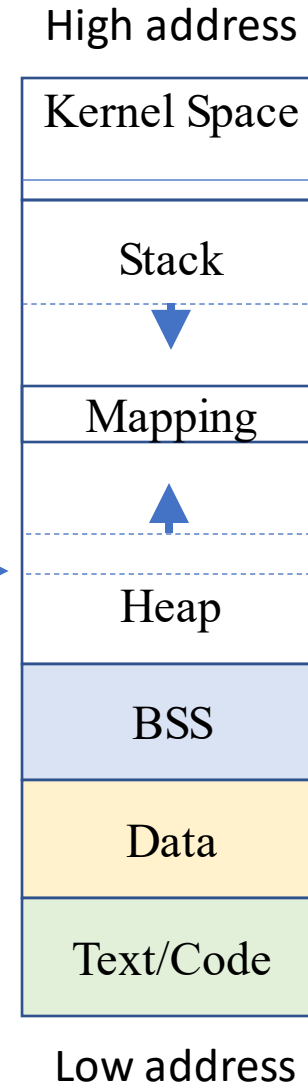
# Program break: The syscall brk

- Sbrk adds a size to the end of cur\_brk

```
void *cur_brk = sbrk(0);
```

```
void *old_brk = sbrk(4096);
```

Program break →



0x8001000: increase 0x8000000 by 4K  
**old\_brk**, cur\_brk: 0x8000000

# Program break: The syscall brk

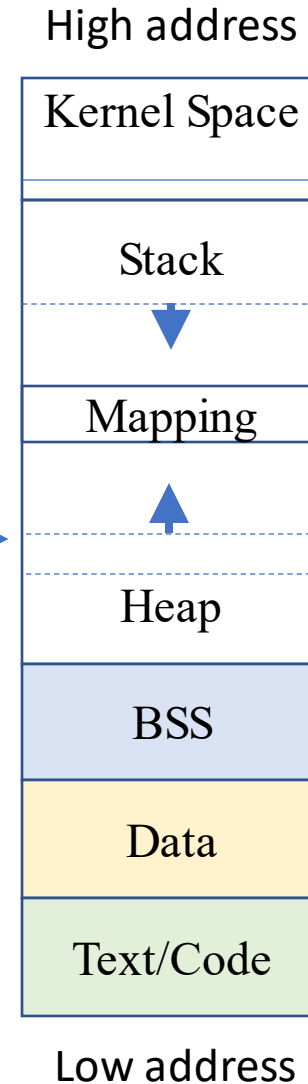
- Sbrk adds a size to the end of cur\_brk

```
void *cur_brk = sbrk(0);
```

```
void *old_brk = sbrk(4096);
```

```
void *new_brk = sbrk(0);
```

Program break →



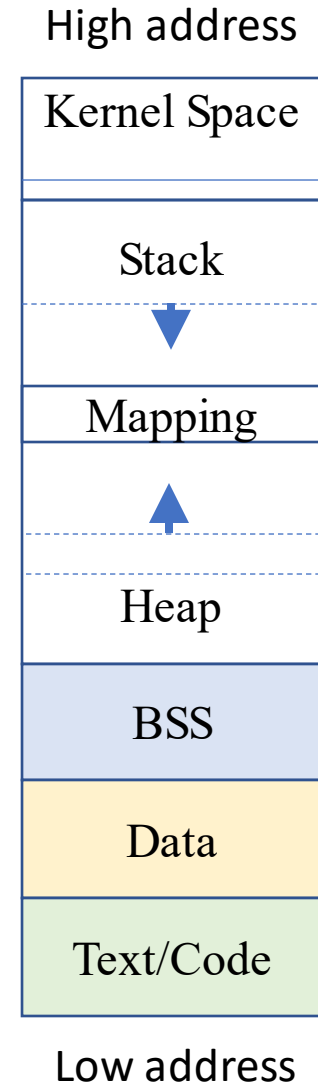
**new\_brk:** 0x8001000

old\_brk, cur\_brk: 0x8000000

# Program break: The syscall brk

---

- brk defines the absolute value for heap's end





# sbrk on XV6

---

The ***sys\_sbrk()*** in ***sysproc.c*** is the XV-6 implementation for sbrk.

```
int
sys_sbrk(void)
{
    int addr;
    int n;

    if(argint(0, &n) < 0)
        return -1;
    addr = myproc() -> sz;
    if(growproc(n) < 0)
        return -1;
    return addr;
}
```

# sbrk on XV6

---

The ***sys\_sbrk()*** in ***sysproc.c*** is the XV-6 implementation for sbrk.

```
int
sys_sbrk(void)
{
    int addr;
    int n;

    if(argint(0, &n) < 0)
        return -1;
    addr = myproc() -> sz; ← Get the current heap size
    if(growproc(n) < 0)
        return -1;
    return addr;
}
```

# sbrk on XV6

---

The ***sys\_sbrk()*** in ***sysproc.c*** is the XV-6 implementation for sbrk.

```
int
sys_sbrk(void)
{
    int addr;
    int n;

    if(argint(0, &n) < 0)
        return -1;
    addr = myproc() -> sz; ← Get the current heapsize
    if(growproc(n) < 0) ← Increase heapsize by n
        return -1;
    return addr;
}
```

# growproc on XV6

---

The ***growproc()*** is in ***proc.c***:

```
int
growproc(int n)
{
    uint sz;
    struct proc *curproc = myproc();

    sz = curproc->sz;
    if(n > 0){
        if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    } else if(n < 0){
        if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    }
    curproc->sz = sz;
    switchuvm(curproc);
    return 0;
}
```

# growproc on XV6

---

The ***growproc()*** is in ***proc.c***:

```
int
growproc(int n)
{
    uint sz;
    struct proc *curproc = myproc();

    sz = curproc->sz;
    if(n > 0){
        if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    } else if(n < 0){
        if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    }
    curproc->sz = sz;
    switchuvm(curproc);
    return 0;
}
```

*Allocates physical page,  
updates page table*

# growproc on XV6

The ***growproc()*** is in ***proc.c***:

```
int
growproc(int n)
{
    uint sz;
    struct proc *curproc = myproc();

    sz = curproc->sz;
    if(n > 0){
        if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    } else if(n < 0){
        if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
            return -1;
    }
    curproc->sz = sz;
    switchuvm(curproc);
    return 0;
}
```

*Allocates physical page,  
updates page table*

*Deallocation, updates page  
table, free physical page*

# growproc on XV6

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?



# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6**: Immediately allocate all 100 physical page frames

# allocvm on xv6

The ***allocvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

# allocuvm on xv6

The ***allocuvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE) {
        mem = kalloc();
        if(mem == 0) {
            cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0) {
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}
```

# allocuvm on xv6

The *allocuvm()* is in *vm*

Allocate page tables and physical memory to grow process from *oldsz* to *newsz*. Return *newsz* if succeed, otherwise

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    int
    growproc(int n)
    {
        uint sz;
        struct proc *curproc = myproc();

        sz = curproc->sz;
        if(n > 0){
            if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
                return -1;
        } else if(n < 0){
            if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
                return -1;
        }
        curproc->sz = sz;
        switchuvm(curproc);
        return 0;
    }
}
```

```
memset(mem, 0, PGSIZE);
if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
    cprintf("allocuvm out of memory (2)\n");
    deallocuvm(pgdir, newsz, oldsz);
    kfree(mem);
    return 0;
}
return newsz;
}
```

# allocuvm on xv6

The ***allocuvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

***mappages(pde\_t \*pgdir, void \*va, unit size, unit pa, int perm)***

Creates translations from ***va*** (virtual address) to ***pa*** (physical address) in existing page table ***pgdir***. Returns 0 if successful, -1 if not.

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE) {
        mem = kalloc();
        if(mem == 0) {
            cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0) {
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}
```

← Round the address to the higher multiple of PGSIZE

← Fill a block of memory with a particular value

Virtual address

# allocuvm on xv6

The ***allocuvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

***mappages(pde\_t \*pgdir, void \*va, unit size, unit pa, int perm)***

Creates translations from ***va*** (virtual address) to ***pa*** (physical address) in existing page table ***pgdir***. Returns 0 if successful, -1 if not.

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNSIZE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE) {
        mem = kalloc();
        if(mem == 0) {
            cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0) {
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}
```

← Round the address to the higher multiple of PGSIZE

← Fill a block of memory with a particular value

← Default page size

# allocuvm on xv6

The ***allocuvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

***mappages(pde\_t \*pgdir, void \*va, unit size, unit pa, int perm)***

Creates translations from ***va*** (virtual address) to ***pa*** (physical address) in existing page table ***pgdir***. Returns 0 if successful, -1 if not.

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE) {
        mem = kalloc();
        if(mem == 0) {
            cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0) {
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}
```

← Round the address to the higher multiple of PGSIZE

← Fill a block of memory with a particular value

← Translating virtual address to physical address

# allocuvm on xv6

The ***allocuvm()*** is in ***vm.c***

Allocate page tables and physical memory to grow process from ***oldsz*** to ***newsz***. Return ***newsz*** if succeed, 0 otherwise

***mappages(pde\_t \*pgdir, void \*va, unit size, unit pa, int perm)***

Creates translations from ***va*** (virtual address) to ***pa*** (physical address) in existing page table ***pgdir***. Returns 0 if successful, -1 if not.

```
int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE) {
        mem = kalloc();
        if(mem == 0) {
            cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0) {
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}
```

← Round the address to the higher multiple of PGSIZE

← Fill a block of memory with a particular value

Flags the page as writeable and to be used by programs (otherwise only the kernel can access it).



# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6:** Immediately allocate all 100 physical page frames

Problems?

# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6**: Immediately allocate all 100 physical page frames

**Lab 4: On Demand Allocation**

# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6**: Immediately allocate all 100 physical page frames

**Lab 4: On Demand Allocation**

# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6:** Immediately allocate all 100 physical page frames

**Lab 4: On Demand Allocation**

allocate one physical page frame upon the 1<sup>st</sup> access on that page.

# Physical Memory Allocation

---

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096\*100
- How about physical memory?

**XV6**: Immediately allocate all 100 physical page frames

## Lab 4: On Demand Allocation

allocate one physical page frame upon the 1<sup>st</sup> access on that page.

allocate one physical page frame when page fault happens.

# Page Fault

---

- Page Table: Stores mapping from virtual page to physical page frame

E.g., Virtual Page 0x8000000 -> Physical 0x400000

- Translating a virtual address to physical address:

Virtual address → (TLB →) Page Table → Physical address

**translation lookaside buffer**

# Page Fault

---

- Translating virtual address 0x8000005:
  1. Get its page-start-address 0x8000000, and **offset**-in-page **5**.
  2. Search (TLB &) Page Table to find the mapping of 0x8000000
  3. If found, e.g., 0x8000000->0x4000000:  
then physical address is 0x4000000**5**.

If not  
found, Page Fault

# Page Fault

---

```
char *ptr = (char*) malloc(4096*100);
```

```
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```



# Page Fault

---

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096\*99 + 50] (**inside** 100<sup>th</sup> page):

# Page Fault

---

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096\*99 + 50] (**inside** 100<sup>th</sup> page):

1. Issue Page Fault trap. All traps are handled by trap() in **trap.c**.

# Page Fault

---

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096\*99 + 50] (**inside** 100<sup>th</sup> page):

1. Issue Page Fault trap. All traps are handled by trap() in trap.c.
2. **Handle Page Fault (*Hint: T\_PGFLT, how to find the faulting addr*) in trap():**
  - 1) **Allocate a physical page frame for this 100<sup>th</sup> page**
  - 2) **Update page table**

# Lab 4 – Part 2 Lazy Allocation

---

- In file **trap.c** and method **trap()**

```
void
trap(struct trapframe *tf)
...
    //cases
...
    //PAGEBREAK: 13
    default:
....
```

# Lab 4 – Part 2 Lazy Allocation

---

- In file **trap.c** and method **trap()**

```
void
trap(struct trapframe *tf)
...
case T_PGFLT:
    growproc(4092);
    break;

//PAGEBREAK: 13
default:
....
```

# Lab 4 – Part 1 Eliminate Allocation from sbrk()

- Just increment the process's size (proc->sz) by n and return the old size.
- Delete the call to growproc()

*Comment out*

```
int
sys_sbrk(void)
{
    int addr;
    int n;

    if(argint(0, &n) < 0)
        return -1;
    addr = myproc()->sz;
    if(growproc(n) < 0)
        return -1;
    return addr;
}
```

# Lab 4 – Part 2 Lazy Allocation

---

- In file **trap.c** and method **trap()**

```
void
```

```
trap(struct trapframe *tf)
```

```
...
```

```
case T_PGFLT:
```

```
    // add just one frame. Exactly where the user requests adjusted by page start address.
```

```
    break;
```

```
//PAGEBREAK: 13
```

```
default:
```

```
....
```

# Lab 4 – Part 2 Lazy Allocation

---

- In file **trap.c** and method **trap()**

```
void
```

```
trap(struct trapframe *tf)
```

```
...
```

```
case T_PGFLT:
```

```
    // add just one frame. Exactly where the user requests adjusted by page start address.
```

```
    PGROUNDDOWN(rcr2());
```

```
    // alloc memory, clean and map it in the process page table.
```

```
    break;
```

```
//PAGEBREAK: 13
```

```
default:
```

```
....
```



# Lab 4 – Part 2 Lazy Allocation

---

- Hint: find the virtual address that caused the page fault
  - In trap.c, find the `cprintf` arguments for “pid XX XX: trap XX err X on cpu X eip ...”
- Hint: you can check whether a fault is a page fault by
  - By checking if ***tf->trapno*** is equal to **T\_PGFLT**
- Hint: reference the logic of `allocuvm()` in `vm.c`
- Hint: use `PGROUNDDOWN(va)` to round the faulting virtual address down to a page
- Hint: break or return in order to avoid the `cprintf` and the `proc->killed = 1`
- Hint: call `int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)`
  - Delete the “static” in the declaration of `mappages()` in `vm.c`
  - Declare `mappages()` in `trap.c`

# Lab 4

---

If all goes well, your lazy allocation code should result in “*echo hi*” working.



# CS 1550

Week 10 – Lab 4

Teaching Assistant

Henrique Potter